

Project 2B

CS111

05/10/19

General Comments on Multithreading

- You noticed that time per operation increases with the number of threads, more so for lists.
 - Does this mean that multithreading is inferior ?
 - Why should we bother with multithreading in that case?

General Comments on Multithreading

- You noticed that time per operation increases with the number of threads, more so for lists.
 - Does this mean that multithreading is inferior ?
 - Why should we bother with multithreading in that case?
- Maybe time per operation is not a great metric
 - It doesn't take parallelism into account
 - What would be better?

General Comments on Multithreading

- You noticed that time per operation increases with the number of threads, more so for lists.
 - Does this mean that multithreading is inferior ?
 - Why should we bother with multithreading at all in that case?
- Maybe time per operation is not a great metric
 - It doesn't take parallelism into account
 - What would be better?
- Throughput : the number of operations per second
 - This allows to take parallelism into account
 - This makes more sense when multithreading is involved

On Contention

- Locks on the shared variable force serialization
 - All threads can run in parallel until they reach the part of code where they acquire the lock
 - At that point they have to wait in line : this is contention
- Why is contention worse for list operations ?

On Contention

- Locks on the shared variable force serialization
 - All threads can run in parallel until they reach the part of code where they acquire the lock
 - At that point they have to wait in line : this is contention
- Why is contention worse for list operations ?
 - Inserting in a list requires more work than addition
 - The **proportion** of time spent in the critical section is higher for list operations
- How could we decrease that contention for lists?

On Contention

- Locks on the shared variable force serialization
 - All threads can run in parallel until they reach the part of code where they acquire the lock
 - At that point they have to wait in line : this is contention
- Why is contention worse for list operations ?
 - Inserting in a list requires more work than addition
 - The **proportion** of time spent in the critical section is higher for list operations
- How could we decrease that contention for lists?
 - We could create several 'queues' to decrease serialization
 - This would be done by splitting up the list in sublists
 - We would lock each sublist independently

Synchronization overhead

- You used 2 types of locks to protect critical sections
 - Spin-locks
 - Mutexes
- Both prevented race conditions
 - Do they add any overhead?
 - Is one better than the other?
- What is the overhead for spin-locks?

Synchronization overhead

- You used 2 types of locks to protect critical sections
 - Spin-locks
 - Mutexes
- Both prevented race conditions
 - Do they add any overhead?
 - Is one better than the other?
- What is the overhead for spin-locks?
 - CPU cycles are used as the thread tries to acquire the lock
 - How can we measure that overhead?

Synchronization overhead

- You used 2 types of locks to protect critical sections
 - Spin-locks
 - Mutexes
- Both prevented race conditions
 - Do they add any overhead?
 - Is one better than the other?
- What is the overhead for spin-locks?
 - CPU cycles are used as the thread tries to acquire the lock
 - How can we measure that overhead?
- What is the overhead for mutexes?
 - Context switches when the thread can't acquire the lock
 - How can we measure that overhead?

Some additions in this project

- `lab2_list.c` -> `--lists` option
 - Accept 'number of lists' as a parameter
 - Initialize several shared lists, each locked independently
 - Decreases contention !
- Execution profiling report. Evaluate:
 - How much time was spent acquiring mutexes?
 - How many cycles were spent acquiring spin-locks?
 - Add an 8th column of data to your `.csv`
- New graphs
 - Using throughput as a metric
- `gperftools` -> multi-threaded application friendly
 - Heap checker
 - Heap profiler
 - Cpu-profiler -> what we'll use to measure spin-lock overhead

How does the cpu profiler work?

- It takes a snapshot of the CPU every millisecond
 - The interval time between snapshots is user controllable
- At each snapshot, it records what function is running
 - This is accumulated over the process runtime
- After the process has run, we can look at the log:
 - If 100 snapshots were taken in function X, we can assume that function ran for 100 ms (which is an approximation ...)
 - In the case of spin locks, what function do we want to monitor?

How does the cpu profiler work?

- It takes a snapshot of the CPU every millisecond
 - The interval time between snapshots is user controllable
- At each snapshot, it records what function is running
 - This is accumulated over the process runtime
- After the process has run, we can look at the log:
 - If 100 snapshots were taken in function X, we can assume that function ran for 100 ms (which is an approximation ...)
 - In the case of spin locks, what function do we want to monitor?
 - `Test_and_Set()`
- The profiler will tell us how many CPU cycles we used to acquire locks
 - This is a good evaluation metric of the overhead

gperftools - Install

- Releases found :
<https://github.com/gperftools/gperftools/releases>
- Unpack and cd to directory
- `./configure --prefix=<path>`
- `make`
- `make install`
- `(make clean)`

Rmk: You do not need root permission to do this

Rmk: You also need to install `gv` to access graphical output, this is optional

gperftools - Use

The pprof command will allow you to analyse the profiling result.

To turn on CPU profiling, you have two options:

1. Define environment variable CPUPROFILE to the filename to dump the profile to (in your makefile)
Rmk: delete that file before each run
2. Bracket the code you want profiled using ProfilerStart() and ProfilerStop(). ProfilerStart() takes the profile filename as argument

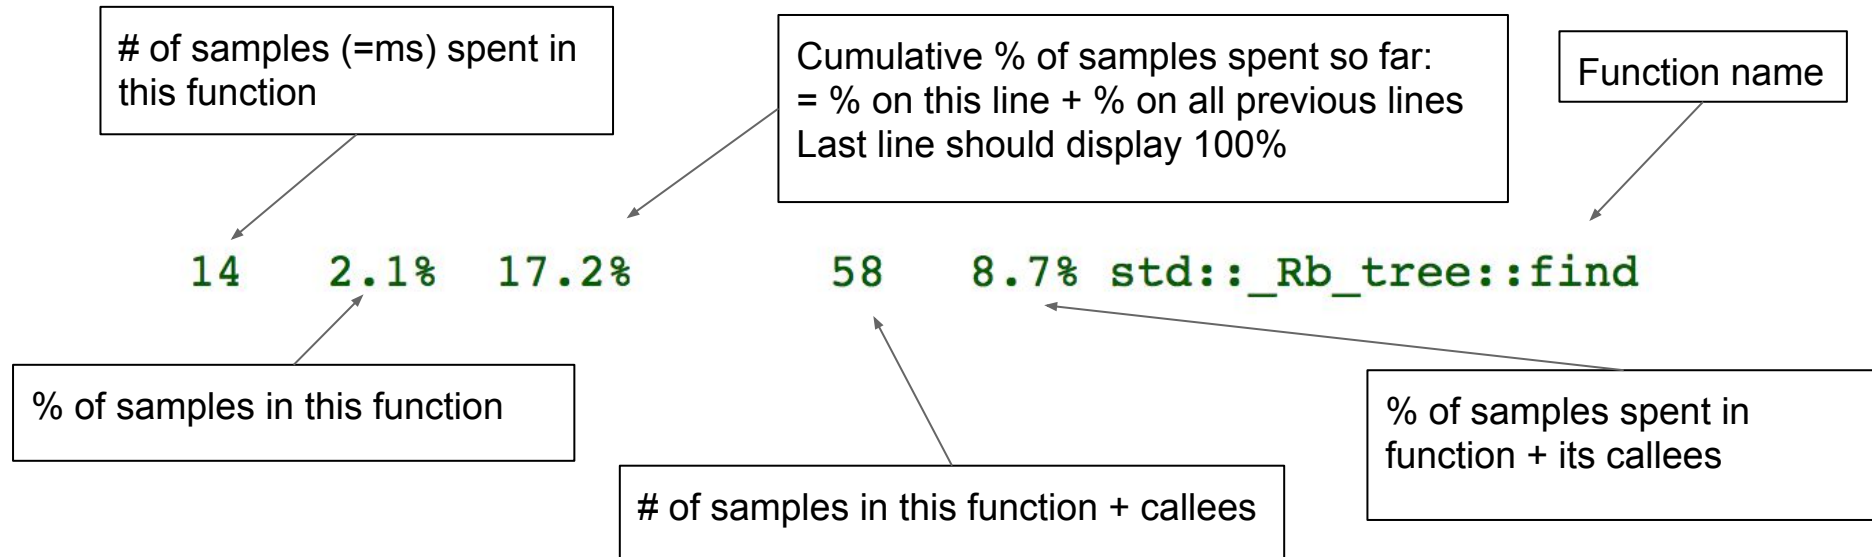
Updating your makefile : profile option

1. Link the library: `path_to_lib` = absolute path where you installed
`LD_PRELOAD = <path_to_lib>` (should be in `/usr/lib(64)`)
2. Link file to record dump
`CPUPROFILE = ./raw.gperf`
3. pprof with `--text`
`pprof --text <executable> <dump file> profile.out`
4. pprof with `--list`
`pprof --list=<function> <exec> <dump> >> profile.out`

Rmk: You also have to compile with the `-lprofile` flag

Analyzing the output

- [For more information](#)
- You **need** to use the text output, but if you first want to use the graphical output to make things clearer you can
- Text output will produce lines that look like:



Gperftools and mutexes

- How can we use gperftools to measure the overhead incurred by mutexes?

Gperftools and mutexes

- How can we use gperftools to measure the overhead incurred by mutexes?
 - We can't!
 - When a mutex can't be acquired, the thread goes to sleep
 - The overhead is a context switch, and time spent sleeping
 - This can't be measured by counting CPU cycles
- What can we do instead?

Gperftools and mutexes

- How can we use gperftools to measure the overhead incurred by mutexes?
 - We can't!
 - When a mutex can't be acquired, the thread goes to sleep
 - The overhead is a context switch, and time spent sleeping
 - This can't be measured by counting CPU cycles
- What can we do instead?
 - We can directly time how long it takes for us to acquire a mutex

Timing Mutex Waits

- Computing wait-for-lock time:
 - Measure time before and after getting the lock
 - Add up all wait time for all threads
 - Divide by number of operations
 - Output Statistics for the run
- If you are not locking -> should get 0
- For example:
 - Allocate an array of timers (# threads)
 - `clock_gettime`, `pthread_mutex_lock([sublist])`, `clock_gettime`
 - add time to `timer[threadnum]`

Addressing the Underlying Problem

- These steps will allow you to prove the original intuition: degradation is a result of increased contention
- To decrease contention, we can split the list
 - Add a lists option to specify the number of chunks
 - Select which list to insert node in
 - Get list length
 - Delete inserted nodes
 - exit

Inserting into different sublists

- There is no requirement of total ordering
 - sublist 1 must be ordered, sublist 2 must be ordered
 - But there is no order requirement **between** sublists
- You are still generating the same number of keys
 - They will be assigned to a single sublist
- To select which list : use a hash function
 - It doesn't matter what hash function you use
 - The goal is to map $\text{threads} * \text{iterations} \rightarrow \# \text{lists}$

Inserting into different sublists

- There is no requirement of total ordering
 - sublist 1 must be ordered, sublist 2 must be ordered
 - But there is no order requirement **between** sublists
- You are still generating the same number of keys
 - They will be assigned to a single sublist
- To select which list : use a hash function
 - It doesn't matter what hash function you use
 - The goal is to map $\text{threads} * \text{iterations} \rightarrow \# \text{lists}$
 - Could be as simple as a modulo operation ($\text{key_value} \% \text{lists}$)

Two sources of efficiency

- Splitting lists increases efficiency by decreasing contention: serialization isn't as bad
- What is the second effect?

Two sources of efficiency

- Splitting lists increases efficiency by decreasing contention: serialization isn't as bad
- What is the second effect?
 - The sublists will be shorter!
 - Therefore insertion operations will be faster
 - The **proportion** of time spent in the critical section will decrease
 - All threads will, on average, wait a shorter amount of time to acquire a lock
 - Fewer cycles/time wasted to acquire a spin-lock/mutex
- These cycles/time can be spent doing useful operations in the critical section