

# Project 3A

## CS111 Discussion 1C

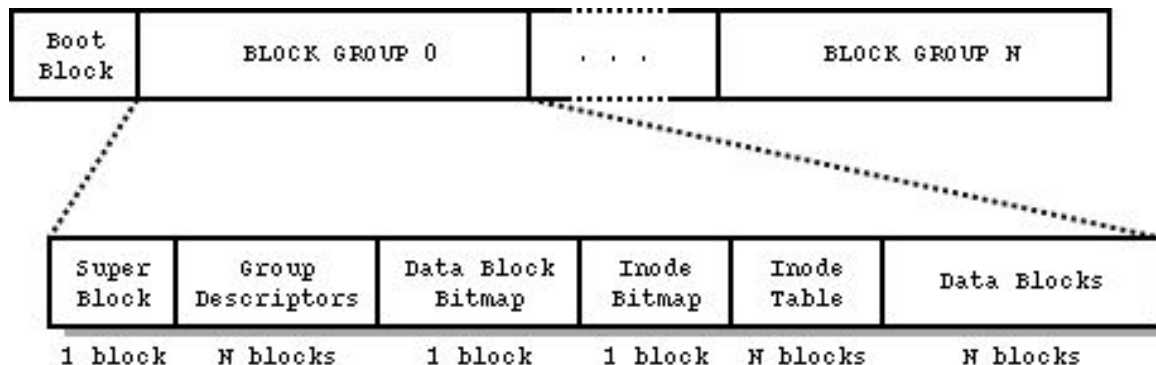
Note: You can form groups of 2 for this project

# EXT2 filesystem

- You will be provided with a filesystem image (EXT2)
- You will explore it
  - You will mount it as a file, which means that:
    - You can read a specific number of byte at a specific offset
    - You will have to repeat these operations for all content
    - The problem becomes calculating the correct offset
- You will output a summary to stdout, describing:
  - The Superblock
  - Groups
  - Free Lists
  - Inodes
  - Indirect Blocks
  - Directories
- This involves a lot of simple code, **if** you have a clear image of what a filesystem is

# EXT2 Filesystem

Read [this](#) carefully. Everything you need is in there.



Where to start? We want to find correct offsets.  
The superblock! We know where it is

Note: File offsets are not addresses in memory! They can't be dereferenced with \*

Note: A function that takes in a block number and outputs its offset would be useful

# Superblock

- Contains all information about the filesystem:
  - # inodes
  - # blocks
  - How many blocks/inodes are free
  - # blocks/inodes per group
  - $\log(\text{block size})$
- Is located at an offset of 1024 bytes from the start of the device
  - This is the only data that we know how to localize, it is the starting point
  - We know the superblock is 1 block long
- You will need to refer to superblock data several times: make it accessible! (save it in a structure)

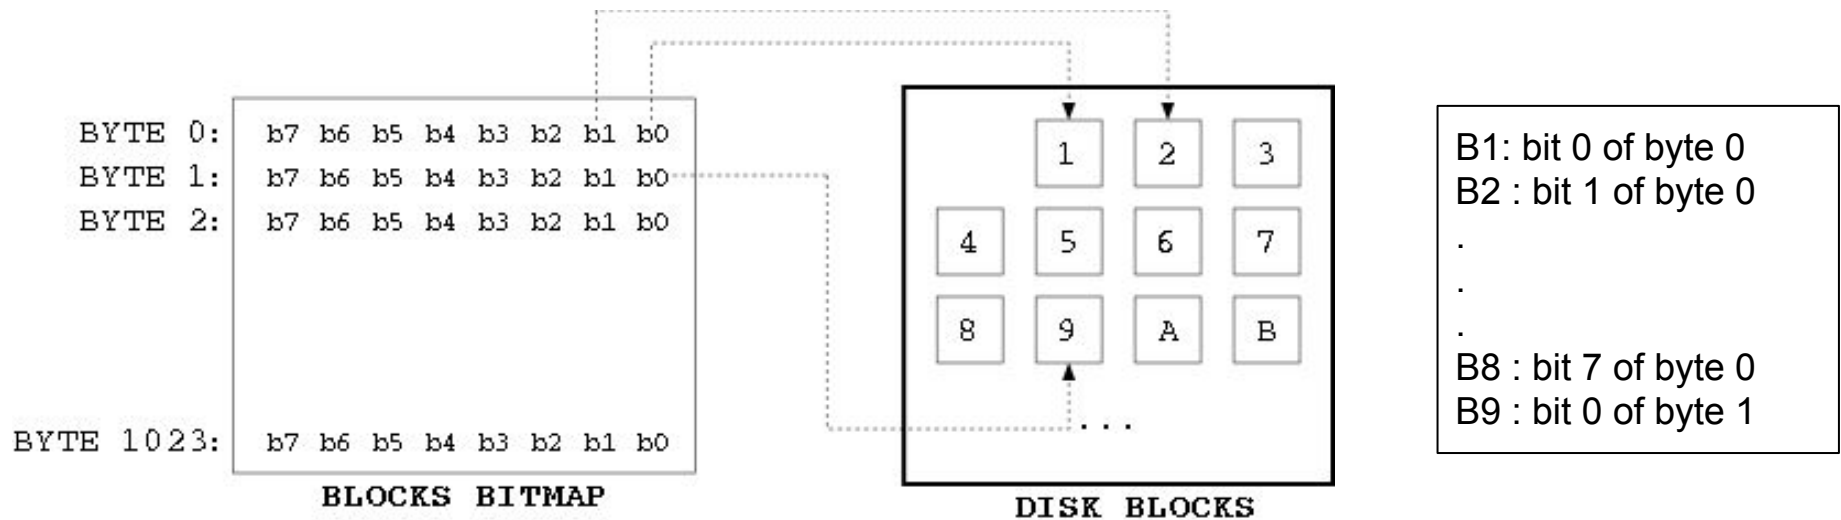
# Block Group Descriptor Table

- Located at block 2, so at  $\text{block\_size} + 1024$ 
  - Block size is in the superblock
- It is an array of block group descriptors. Each of these block group descriptors provide (for 1 group):
  - The location of the inode bitmap/table
  - The location of the block bitmap
  - # of free blocks/inode
- This table can be stored on multiple blocks
  - How many? -> Superblock

Superblock offset

# Block Bitmap

- Where is it located?
  - Bg\_block\_bitmap of its block group descriptor
- What does it represent?
  - Current state of a block within that block group (1 = used / 0=free)
- How is the information represented?



# Inode Bitmap

- Where is it located?
  - Bg\_inode\_bitmap in group descriptor
- What does it represent?
  - Current state of an inode within the inode table (1=used)
- Works similarly to a Block Bitmap

# Inode Table

- Where is it located?
  - `bg_inode_table` in the group descriptor
- It is an array of inodes. Each of those represent a single file (a directory, socket, buffer, symbolic link, regular file):
  - size and owner
  - access/modification/creation times (seconds since 1/1/70)
  - the number of blocks containing the data of the inode
  - which blocks contain the data of the file pointed to by the inode
  - Note that there is no filename in the inode
- The first few entries of the inode table are reserved



# Locating an inode


- They are numerically ordered
  - Inode number  $\rightarrow$  index in the inode table
- The size of the inode table is fixed at format time
  - There is a cap on the number of entries
  - The size is (# of inodes \* inode size) (cf superblock)
- Therefore the local inode index (relative to the current block group) is :
  - $(\text{inode\_idx} - 1) \% \text{inodes\_per\_group}$

# Inode pointer structure

- The inode points to the blocks that contain its data
- The block numbers are stored in a table
  - The first 12 are 'direct' blocks
  - The 13th block is the 'indirect block'
  - The 14th block is the 'doubly indirect block'
  - The 15th block is the 'triply indirect block'
  - If a 0 is encountered, all subsequent block pointers should be 0 (no other blocks defined)
- The indirect block structure is used to be able to reference files that would need more than 15 blocks with a single (fixed size inode)

# Inode pointer structure

A picture is probably helpful:



# Debugfs

- Download the image and mount it
  - This requires sudo: work on your machine
- You can explore it using debugfs(8)
  - You will need to research this tool to properly interpret values

```

debugfs: stat file1
Inode: 2790782   Type: regular   Mode:  0600   Flags: 0x0   Generation: 46520506
User:  2605     Group:  2601     Size: 14
File ACL: 0     Directory ACL: 0
Links: 1      Blockcount: 8
Fragment:  Address: 0      Number: 0      Size: 0
ctime: 0x3be712ea -- Mon Nov  5 15:30:02 2001
atime: 0x3be712ea -- Mon Nov  5 15:30:02 2001
mtime: 0x3be712ea -- Mon Nov  5 15:30:02 2001
BLOCKS:
5603924
TOTAL: 1
  
```

Dump content of inode

Access Control List

Inode change/access/modification time

# Debugfs

- Other useful commands include:
  - `bd` : block dump
  - `testi <inode>` : test if inode is marked allocated in the inode bitmap
  - `testb <block>` : test if block is marked allocated in the block bitmap
- You can use this tool while writing your program to verify that you are exploring the filesystem correctly

# General code structure

- Find/analyze/report superblock
- Copy info to common structure
  - Make the superblock data easily available to you
- Read group descriptions
  - How many blocks per group?
  - How many inodes per group?
- Find and report both bitmaps
- Analyze all allocated inodes
  - Is it a directory? A regular file?
  - What is the level of indirection?

# Some useful notes

- Number of blocks per group?

# Some useful notes

- Number of blocks per group?
  - You can find this in the superblock, except for the last group
- What is the block size?



# Some useful notes

- Number of blocks per group?
  - You can find this in the superblock, except for the last group
- What is the block size?
  - $1024 \ll \text{superblock.s\_log\_block\_size}$
- Is the file a regular file or a directory?

# Some useful notes

- Number of blocks per group?
  - You can find this in the superblock, except for the last group
- What is the block size?
  - $1024 \ll \text{superblock.s\_log\_block\_size}$
- Is the file a regular file or a directory?
  - Answer is in the first field of the inode : the 'mode'
- Write a function that takes in a block number and returns the absolute address (offset) of that block
- Write a function that reads a given number of blocks at a given offset

# Iterating through inodes

- What is the level of indirection?
  - For a 1KiB block size:
    - Block 1 to 12 are direct
    - Block 13 to 268 are indirect
    - How many double indirect blocks?

# Iterating through inodes

- What is the level of indirection?
  - For a 1KiB block size:
    - Block 1 to 12 are direct
    - Block 13 to 268 are indirect
    - How many double indirect blocks?  $(256)*(256) = 65536$
    - How many triple indirect blocks?

# Iterating through inodes

- What is the level of indirection?
  - For a 1KiB block size:
    - Block 1 to 12 are direct
    - Block 13 to 268 are indirect
    - How many double indirect blocks?  $(256)*(256) = 65536$
    - How many triple indirect blocks?  $(256)*(256)*(256) = 16777216$
- How does your offset change when you go through it?
  - Direct blocks -> blocksize
  - Indirect blocks?

# Iterating through inodes

- What is the level of indirection?
  - For a 1KiB block size:
    - Block 1 to 12 are direct
    - Block 13 to 268 are indirect
    - How many double indirect blocks?  $(256) * (256) = 65536$
    - How many triple indirect blocks?  $(256) * (256) * (256) = 16777216$
- How does your offset change when you go through it?
  - Direct blocks -> blocksize
  - Indirect blocks? Multiply by blocksize per indirection level

# Directory Structure

- Stored as datablocks, pointed to by an inode
- Of type EXT2\_S\_IFDIR
- Stored as a linked list
- How do you find a file?

# A word on the trivial image

- You are provided with a small filesystem and the correct analysis output for that filesystem
- You are expected to be able to reproduce 'trivial.csv' from the given filesystem

## BUT

- Doing so won't guarantee that your program is 100% correct, this tests basic functionality
- Also note that any access will modify the inodes!
  - This means that you won't be able to refer to the provided .csv
  - Use a fresh copy / mount in read only



# FAQ

- How do I mount the image on a mac?  
`hdiutil attach -readonly -nomount ./trivial.img`
- Can I assume that there will only be a single group in the filesystem used for grading?  
Yes!